

Uncertainty Expression Gate: Complete Implementation

For ICLR Workshop: Principled Design for Trustworthy AI

This repository contains complete, working code for all 5 experiments demonstrating that instruction-tuned language models contain a policy-responsive gate that controls whether internal uncertainty is expressed as abstention.

🎯 Key Contribution

We prove that abstention is governed by a late-layer "expression gate" that:

1. Responds to instruction regime (not just question properties)
2. Operates independently of internal uncertainty
3. Is causally manipulable via targeted interventions
4. Enables tuning trustworthiness tradeoffs

📁 File Structure

```
.  
├── core_utils.py      # Model wrapper & utilities  
├── data_preparation.py    # Dataset creation  
├── experiment1_behavior_belief.py  # Exp 1: Policy-responsiveness  
├── experiment2_localization.py   # Exp 2: Layer localization  
├── experiment3_4_steering_independence.py # Exp 3-4: Steering & key proof  
├── experiment5_trustworthiness.py  # Exp 5: Application + master runner  
├── data/                # Generated datasets (auto-created)  
└── results/             # Output figures & CSVs (auto-created)
```

🚀 Quick Start

1. Installation

```
bash
```

```
# Create environment
conda create -n mech_interp python=3.10
conda activate mech_interp

# Install dependencies
pip install torch transformers numpy pandas matplotlib seaborn scikit-learn tqdm
```

2. Prepare Data

```
bash
python data_preparation.py
```

This creates:

- `data/dataset_ambiguous.json` (30 ambiguous questions)
- `data/dataset_clearly_answerable.json` (10 factual questions)
- `data/dataset_clearly_unanswerable.json` (10 unanswerable questions)
- `data/dataset_squad.json` (synthetic SQuAD-style examples)

Optional: Replace with real SQuAD v2.0 data by downloading from <https://rajpurkar.github.io/SQuAD-explorer/> and updating `load_squad_data()` in `data_preparation.py`.

3. Run Experiments

Option A: Quick Test (recommended first)

```
bash
python experiments5_trustworthiness.py
# This runs quick_test() with minimal data (~15 min on GPU)
```

Option B: Run Individual Experiments

```
bash
```

```

# Experiment 1: ~30 min
python experiment1_behavior_belief.py

# Experiment 2: ~45 min (tests all layers)
python experiment2_localization.py

# Experiments 3 & 4: ~1 hour
python -c "from experiment3_4_steering_independence import main; main()"

# Experiment 5: ~20 min
python -c "from experiment5_trustworthiness import main; main()"

```

Option C: Run All Experiments (full pipeline)

```

bash

python -c "from experiment5_trustworthiness import run_all_experiments; run_all_experiments(quick_test=False)"

# Full run: ~3-4 hours on GPU

```

4. View Results

All outputs saved to `results/`:

- `exp1_behavior_belief_dissociation.png` - Instruction regime effects
- `exp2_localization_analysis.png` - Layer-wise localization
- `exp3_steering_analysis.png` - Steering control
- `exp4_gate_independence.png` - ★ KEY FIGURE: Gate independence proof
- `exp5_trustworthiness.png` - Risk-coverage tradeoff
- `master_summary.json` - All quantitative results

Experiments Overview

Experiment 1: Behavior ≠ Belief

Shows: Instruction regime dramatically changes abstention behavior while internal uncertainty stays constant.

Key metric: Behavior gap (cautious - confident) vs Belief gap (entropy difference)

Experiment 2: Gate Localization

Shows: Abstention decision is causally determined in a late-layer band (typically layers 17+).

Method: Activation patching from answerable → unanswerable examples

Experiment 3: Low-Dimensional Control

Shows: A single learned direction can steer abstention behavior reliably.

Method: Compute answerability direction via mean difference, test steering

Experiment 4: Gate Independence ★ CRITICAL

Shows: Gate operates independently of uncertainty - can force high-uncertainty questions to be answered, and low-uncertainty questions to be abstained.

This is your novel contribution!

Experiment 5: Trustworthiness Application

Shows: Intervening on the gate enables tuning the risk-coverage tradeoff.

Method: Selective accuracy evaluation across steering strengths

🔧 Customization

Using Your Own Data

Replace datasets in `(data_preparation.py)`:

```
python

def create_your_custom_dataset():
    return [
        {
            "question": "Your question here",
            "context": "Optional context",
            "answer": "Ground truth" or None,
            "answerability": "answerable" or "unanswerable"
        },
        # ... more examples
    ]
```

Adjusting Compute Requirements

For faster testing:

```
python
```

```
# In experiment scripts
config.n_force_guess_samples = 10 # Reduce from 20
layer_stride = 2 # Test every other layer in Exp 2
```

For better results:

```
python

config.n_force_guess_samples = 30 # More samples = better entropy estimates
# Test more example pairs in Exp 2, 3, 4
```

Different Models

Change model in `core_utils.py`:

```
python

model_name: str = "Qwen/Qwen2.5-1.5B-Instruct" # Current
# Try:
# "Qwen/Qwen2.5-7B-Instruct"
# "meta-llama/Llama-2-7b-chat-hf"
# etc.
```

Note: You may need to adjust `config.target_layers` based on model depth.

✓ Expected Results

Experiment 1

- Abstention rate should vary 30-60% between cautious and confident regimes
- Internal entropy should vary <10% across regimes
- Behavior/Belief ratio should be >3x

Experiment 2

- Clear concentration of causal effect in late layers (typically last 20-30%)
- Flip rates should exceed 50% in critical layers

Experiment 3

- Flip rates should increase monotonically with epsilon

- Layer ~26 typically most effective (for 28-layer model)

Experiment 4

- Should observe cases where:
 - High uncertainty → forced to answer (gate opens)
 - Low uncertainty → forced to abstain (gate closes)
- This proves gate independence!

Experiment 5

- Risk-coverage curve should show tradeoff
- Higher epsilon → lower coverage but higher accuracy

Troubleshooting

Out of Memory

```
python

# Reduce batch size / samples
config.n_force_guess_samples = 10
# Use smaller model
model_name = "Qwen/Qwen2.5-1.5B-Instruct" # instead of 7B
```

Model Loading Issues

```
bash

# Make sure you have access to model
huggingface-cli login
# Or use different model that doesn't require auth
```

Incorrect Behavior Detection

The `extract_answer()` function uses simple heuristics. For better detection:

```
python
```

```
# Add more uncertainty markers
uncertainty_markers = [
    "UNCERTAIN", "I don't know", "cannot answer",
    "can't determine", "insufficient information",
    # ... add markers specific to your model
]
```

Layer Index Errors

Different model architectures have different layer counts. Check:

```
python
n_layers = model.model.config.num_hidden_layers
config.target_layers = [n_layers-4, n_layers-3, n_layers-2, n_layers-1]
```

Writing the Paper

Key Claims to Make

1. **Abstention is governed by a policy-responsive gate**
 - Evidence: Exp 1 (behavior changes, belief doesn't)
2. **Gate operates in late layers**
 - Evidence: Exp 2 (activation patching localization)
3. **Gate is low-dimensional and manipulable**
 - Evidence: Exp 3 (steering control)
4. **Gate operates independently of uncertainty ★**
 - Evidence: Exp 4 (changes abstention at fixed uncertainty)
5. **Practical trustworthiness implications**
 - Evidence: Exp 5 (risk-coverage tuning)

Figures for Paper

Main text (4 figures max):

1. Fig 1: Exp 1 behavior-belief dissociation
2. Fig 2: Exp 2 localization (patching effects by layer)

3. Fig 3: Exp 4 gate independence  (this is your money figure!)

4. Fig 4: Exp 5 risk-coverage curve

Supplement:

- Exp 3 steering details
- Additional ablations
- Failure cases

One-Sentence Summary

"We identify a late, low-dimensional policy gate in instruction-tuned language models that controls whether internal uncertainty is expressed as abstention, and show that intervening on this gate changes abstention behavior at fixed internal answerability."

Citation

If you use this code, please cite:

```
bibtex
```

```
@inproceedings{yourname2025uncertainty,  
    title={Uncertainty Expression Gate: Mechanistic Control of Abstention in Language Models},  
    author={Your Name},  
    booktitle={ICLR Workshop on Principled Design for Trustworthy AI},  
    year={2025}  
}
```

Support

For questions or issues:

1. Check troubleshooting section above

2. Review code comments

3. Open an issue with:

- Error message
- Which experiment
- Your config settings

Timeline for Jan 30 Deadline

- **Week 1 (Jan 2-8):** Run all experiments, debug issues
- **Week 2 (Jan 9-15):** Analyze results, create figures
- **Week 3 (Jan 16-22):** Write draft, iterate on experiments if needed
- **Week 4 (Jan 23-30):** Finalize paper, submission

Good luck! You've got this! 